


```
#import the libraries
from google.colab import drive
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.feature_selection import mutual_info_classif
from sklearn.feature_selection import SelectKBest
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, precision_recall_curve
import matplotlib.pyplot as plt
import seaborn as sns
```

```
#read the dataset
drive.mount('/content/drive')
df_train = pd.read_csv("/content/drive/MyDrive/archive/fraudTrain.csv")
df_test = pd.read_csv("/content/drive/MyDrive/archive/fraudTest.csv")
```

Mounted at /content/drive

```
#read the dataset
df = pd.concat([df_train, df_test], ignore_index=True, axis=0)
df.head()
```

 Unnamed: 0

	0	trans_date	trans_time	cc_num	merchant	category	amt	first	last	gender	street	...	
0	0	2019-01-01	00:00:18	2703186189652095	fraud_Rippin, Kub and Mann	misc_net	4.97	Jennifer	Banks	F	561 Perry Cove	...	36.0
1	1	2019-01-01	00:00:44	630423337322	fraud_Heller, Gutmann and Zieme	grocery_pos	107.23	Stephanie	Gill	F	43039 Riley Greens Suite 393	...	48.8
2	2	2019-01-01	00:00:51	38859492057661	fraud_Lind-Buckridge	entertainment	220.11	Edward	Sanchez	M	594 White Dale Suite 530	...	42.1
3	3	2019-01-01	00:01:16	3534093764340240	fraud_Kutch, Hermiston and Farrell	gas_transport	45.00	Jeremy	White	M	9443 Cynthia Court Apt. 038	...	46.2
4	4	2019-01-01	00:03:06	375534208663984	fraud_Keeling-Crist	misc_pos	41.96	Tyler	Garcia	M	408 Bradley Rest	...	38.4

5 rows × 23 columns

```
df.describe()
```

	Unnamed: 0	cc_num	amt	zip	lat	lon
count	1.852394e+06	1.852394e+06	1.852394e+06	1.852394e+06	1.852394e+06	1.852394e+06
mean	5.371934e+05	4.173860e+17	7.006357e+01	4.881326e+04	3.853931e+01	-9.022783e+01
std	3.669110e+05	1.309115e+18	1.592540e+02	2.688185e+04	5.071470e+00	1.374789e+01
min	0.000000e+00	6.041621e+10	1.000000e+00	1.257000e+03	2.002710e+01	-1.656723e+01
25%	2.315490e+05	1.800429e+14	9.640000e+00	2.623700e+04	3.466890e+01	-9.679800e+01
50%	4.630980e+05	3.521417e+15	4.745000e+01	4.817400e+04	3.935430e+01	-8.747690e+01
75%	8.335758e+05	4.642255e+15	8.310000e+01	7.204200e+04	4.194040e+01	-8.015800e+01
max	1.296674e+06	4.992346e+18	2.894890e+04	9.992100e+04	6.669330e+01	-6.795030e+01

```
#legit and fraudulent
print(df['is_fraud'].value_counts())
legit = df[df.is_fraud==0]
fraudulent = df[df.is_fraud==1]
print(f"legit data: {legit.shape}")
print(f"fraudulent data: {fraudulent.shape}")
```

```
0    1842743
1      9651
Name: is_fraud, dtype: int64
legit data: (1842743, 23)
fraudlent data: (9651, 23)
```

```
#under sampling
legit = legit.sample(n=9700)
print(f"legit data: {legit.shape}")
print(f"fraudlent data: {fraudlent.shape}")
df=pd.concat([legit,fraudlent],axis=0)
print(df.sort_index())
```

```
legit data: (9700, 23)
fraudlent data: (9651, 23)
Unnamed: 0 trans_date_trans_time cc_num \
176      176  2019-01-01 02:10:12  4740713119940984
271      271  2019-01-01 03:30:49  36078114201167
360      360  2019-01-01 04:42:23  38588538868506
506      506  2019-01-01 06:39:59  3590736522064285
659      659  2019-01-01 08:36:25  3583635130604947
...      ...      ...      ...
1851288  554613  2020-12-31 17:37:10  4265776278887457
1851454  554779  2020-12-31 18:28:41  4223708906367574214
1851793  555118  2020-12-31 20:30:50  213161231269724
1851883  555208  2020-12-31 21:02:06  30030380240193
1852105  555430  2020-12-31 22:15:27  30407675418785
```

```
merchant category amt \
176      fraud_Brown-Greenholt entertainment 5.27
271      fraud_Deckow-O'Conner grocery_pos 145.21
360      fraud_Schaefer, Maggio and Daugherty gas_transport 60.20
506      fraud_Stamm-Rodriguez misc_pos 76.40
659      fraud_Bauch-Raynor grocery_pos 206.85
...      ...      ...      ...
1851288      fraud_Harris Group food_dining 91.13
1851454      fraud_Hauck, Dietrich and Funk kids_pets 49.45
1851793      fraud_Gerhold LLC home 32.20
1851883      fraud_Hyatt, Russel and Gleichner health_fitness 76.44
1852105      fraud_Kris-Padberg shopping_pos 12.88
```

```
first last gender street ... \
176      Heather Hines F 13776 Hicks Plains ...
271      Christopher Horn M 956 Sanchez Highway ...
360      Jacqueline Curry F 3047 Jeff Place ...
506      Kimberly Gonzalez F 72966 Shannon Pass Apt. 391 ...
659      Crystal Gamble F 899 Michele View Suite 960 ...
...      ...      ...      ...
1851288      Christine Best F 68248 Deanna Land ...
1851454      Adam Riddle M 27718 Mason Bypass ...
1851793      Alyssa Morgan F 622 Robin Run Suite 764 ...
1851883      William Jenkins M 50614 Kevin Point ...
1852105      Danielle Evans F 76752 David Lodge Apt. 064 ...
```

```
lat long city_pop job dob \
176  41.1901 -74.0436 9993 Pensions consultant 1962-10-16
271  37.2692 -82.9161 798 Facilities manager 1926-06-26
360  30.1886 -103.2214 498 Lexicographer 1990-11-23
506  34.5091 -92.4828 4074 Scientist, audiological 1975-12-20
659  40.0369 -75.0664 1526206 Structural engineer 1985-01-01
...      ...      ...      ...
1851288  35.2087 -92.2123 969 Physicist, medical 1954-01-05
1851454  39.0965 -84.6431 177 Exhibition designer 1974-05-30
1851793  34.0480 -85.9246 67082 Physiological scientist 1963-02-09
1851883  30.2816 -99.2410 2395 Pharmacist, community 1993-11-17
1852105  42.1939 -76.7361 520 Psychotherapist 1991-10-13
```

```
trans_num unix_time merch_lat merch_long \
176  9f55bd65f64193f1b1a46a99c93e7844 1325383812 41.788775 -74.471154
271  50fd631627ff1c488113e8c9249ab801 1325388649 37.693205 -82.671226
360  9da7e99fc8147da68b5e322d7c63c096 1325392943 29.335557 -103.852785
```

```
#statistical measures
print(df.describe())
print(df.groupby('is_fraud').mean())
```

```
Unnamed: 0 cc_num amt zip lat \
count 1.935100e+04 1.935100e+04 19351.000000 19351.000000 19351.000000
mean 5.386690e+05 3.975616e+17 298.099432 48464.971423 38.663691
std 3.792391e+05 1.277992e+18 375.506985 27079.031702 5.111601
min 1.760000e+02 6.041621e+10 1.000000 1257.000000 20.027100
25% 2.150450e+05 1.800365e+14 20.170000 25213.000000 34.847000
50% 4.593270e+05 3.520550e+15 88.850000 47863.000000 39.433600
75% 8.667740e+05 4.633065e+15 465.165000 71762.000000 42.074000
max 1.296623e+06 4.992346e+18 8517.380000 99921.000000 66.693300

long city_pop unix_time merch_lat merch_long \
```

```

count    19351.000000    1.935100e+04    1.935100e+04    19351.000000    19351.000000
mean      -90.209527     8.794817e+04    1.357116e+09    38.661089     -90.210209
std       14.130837     2.979177e+05    1.817924e+07     5.145979     14.144194
min      -165.672300     2.300000e+01    1.325384e+09    19.121455    -166.562839
25%      -96.786900     7.540000e+02    1.341615e+09    34.930664    -96.841426
50%      -87.349000     2.526000e+03    1.356332e+09    39.473030    -87.281755
75%      -80.065200     1.940800e+04    1.372604e+09    42.016427    -80.102078
max      -67.950300     2.906700e+06    1.388528e+09    67.510267    -66.960745

```

```
is_fraud
```

```

count    19351.000000
mean       0.498734
std        0.500011
min        0.000000
25%        0.000000
50%        0.000000
75%        1.000000
max        1.000000

```

```

Unnamed: 0      cc_num      amt      zip      lat \
is_fraud
0      537494.653918    3.980812e+17    66.712249    48927.953814    38.584969
1      539849.247228    3.970393e+17    530.661412    47999.638379    38.742813

```

```

long      city_pop      unix_time      merch_lat      merch_long
is_fraud
0      -90.384436    85908.282784    1.358900e+09    38.587590    -90.381628
1      -90.033730    89998.422961    1.355323e+09    38.734962    -90.037919

```

```

<ipython-input-7-c5a531823b49>:3: FutureWarning: The default value of numeric_only in DataFrameGroupBy.mean is deprecated. In a futu
print(df.groupby('is_fraud').mean())

```

```

#print number of unique categories for every attribute
print(f"Total number of rows: {df.shape[0]}")
for col in df.columns:
    print(f'{col}: {df[col].nunique()}')

```

```

Total number of rows: 19351
Unnamed: 0: 19300
trans_date_trans_time: 19345
cc_num: 999
merchant: 693
category: 14
amt: 13994
first: 355
last: 486
gender: 2
street: 999
city: 906
state: 51
zip: 985
lat: 983
long: 983
city_pop: 891
job: 497
dob: 984
trans_num: 19351
unix_time: 19345
merch_lat: 19338
merch_long: 19345
is_fraud: 2

```

```

#dropping columns and check for null values
df.drop(labels=['Unnamed: 0'], axis=1, inplace=True)
print(f"Null values: {df.isnull().values.any()}")
print(df.columns)

```

```

Null values: False
Index(['trans_date_trans_time', 'cc_num', 'merchant', 'category', 'amt',
      'first', 'last', 'gender', 'street', 'city', 'state', 'zip', 'lat',
      'long', 'city_pop', 'job', 'dob', 'trans_num', 'unix_time', 'merch_lat',
      'merch_long', 'is_fraud'],
      dtype='object')

```

```

#calculate entropy for target variable
target = "is_fraud"
probabilities = df[target].value_counts(normalize=True)/len(df)
entropy = -np.sum(probabilities*np.log2(probabilities))
print(f"{target}: {entropy}")

```

```
is_fraud: 0.0007875621867953986
```

```

#change string labels to numerical
label_encoder = LabelEncoder()
stringTypeColumns = ['trans_date_trans_time','cc_num', 'merchant', 'category', 'first', 'last', 'gender', 'street', 'city', 'state', 'zip']
for column in stringTypeColumns:
    df[column] = label_encoder.fit_transform(df[column])

#calculate information gain for the features
X=df.drop('is_fraud', axis=1)
y=df['is_fraud']
information_gain = mutual_info_classif(X, y, discrete_features='auto', random_state=1)
information_gain

array([0.25748812, 0.08298698, 0.09281903, 0.08651023, 0.4525601 ,
       0.03049561, 0.03859248, 0.00135004, 0.07616792, 0.07114658,
       0.00488369, 0.07923783, 0.07601854, 0.07323542, 0.06354338,
       0.03173214, 0.0775402 , 0.00234236, 0.24890171, 0.00185617,
       0.0011095 ])

#print the information gain
info_gain = pd.Series(information_gain)
info_gain.index = X.columns
info_gain.sort_values(ascending= False)

amt                0.452560
trans_date_trans_time 0.257488
unix_time          0.248902
merchant           0.092819
category           0.086510
cc_num             0.082987
zip                0.079238
dob                0.077540
street             0.076168
lat                0.076019
long               0.073235
city               0.071147
city_pop           0.063543
last               0.038592
job                0.031732
first              0.030496
state              0.004884
trans_num          0.002342
merch_lat          0.001856
gender             0.001350
merch_long         0.001109
dtype: float64

#plot the information gain graph
ax = info_gain.sort_values(ascending = True).plot(kind='barh', color='blue', figsize=(16, 6))

plt.xlabel('Information Gains')
plt.ylabel('Categorical Variables')
plt.title('Information Gains vs Categorical Variables')

plt.show()

```

Information Gains vs Categorical Variables

```
#selecting features and create new dataframe
select_cols = SelectKBest(mutual_info_classif, k=10)
select_cols.fit_transform(X,y)
selected_indices = select_cols.get_support(indices=True)
selected_indices = list(selected_indices) + [df.columns.get_loc('is_fraud')]
new_df = df.iloc[:, selected_indices]
print(new_df.sort_index())
```

	trans_date_trans_time	cc_num	merchant	category	amt	street	\
176	0	767	81	0	5.27	134	
271	1	219	130	4	145.21	963	
360	2	238	548	2	60.20	314	
506	3	627	592	9	76.40	738	
659	4	616	29	4	206.85	908	
...	
1851288	19340	688	230	1	91.13	692	
1851454	19341	924	234	7	49.45	279	
1851793	19342	295	188	6	32.20	636	
1851883	19343	157	274	5	76.44	517	
1852105	19344	198	341	12	12.88	785	

	zip	lat	long	unix_time	is_fraud
176	70	41.1901	-74.0436	1325383812	0
271	427	37.2692	-82.9161	1325388649	0
360	844	30.1886	-103.2214	1325392943	0
506	745	34.5091	-92.4828	1325399999	0
659	182	40.0369	-75.0664	1325406985	0
...
1851288	747	35.2087	-92.2123	1388511430	0
1851454	458	39.0965	-84.6431	1388514521	0
1851793	360	34.0480	-85.9246	1388521850	0
1851883	831	30.2816	-99.2410	1388523726	0
1852105	130	42.1939	-76.7361	1388528127	0

[19351 rows x 11 columns]

```
#print new dataframe statistical measures
print(f"Total number of rows: {new_df.shape[0]}")
for col in new_df.columns:
    print(f'{col}: {new_df[col].nunique()}')
print(new_df.describe())
print(new_df.groupby('is_fraud').mean())
```

```
Total number of rows: 19351
trans_date_trans_time: 19345
cc_num: 999
merchant: 693
category: 14
amt: 13994
street: 999
zip: 985
lat: 983
long: 983
unix_time: 19345
is_fraud: 2
```

	trans_date_trans_time	cc_num	merchant	category	\
count	19351.000000	19351.000000	19351.000000	19351.000000	
mean	9672.016537	497.293266	341.036949	6.734277	
std	5584.208423	286.417764	197.287791	3.870311	
min	0.000000	0.000000	0.000000	0.000000	
25%	4836.500000	249.000000	173.000000	4.000000	
50%	9672.000000	493.000000	345.000000	7.000000	
75%	14506.500000	744.000000	504.000000	11.000000	
max	19344.000000	998.000000	692.000000	13.000000	

	amt	street	zip	lat	long	\
count	19351.000000	19351.000000	19351.000000	19351.000000	19351.000000	
mean	298.099432	497.363030	491.162421	38.663691	-90.209527	
std	375.506985	286.954735	285.970290	5.111601	14.130837	
min	1.000000	0.000000	0.000000	20.027100	-165.672300	
25%	20.170000	253.000000	242.000000	34.847000	-96.786900	
50%	88.850000	495.000000	493.000000	39.433600	-87.349000	
75%	465.165000	743.000000	741.000000	42.074000	-80.065200	
max	8517.380000	998.000000	984.000000	66.693300	-67.950300	

	unix_time	is_fraud
count	1.935100e+04	19351.000000
mean	1.357116e+09	0.498734
std	1.817924e+07	0.500011
min	1.325384e+09	0.000000
25%	1.341615e+09	0.000000
50%	1.356332e+09	0.000000
75%	1.372604e+09	1.000000
max	1.388528e+09	1.000000

	trans_date_trans_time	cc_num	merchant	category	amt	\
--	-----------------------	--------	----------	----------	-----	---

```

is_fraud
0      10217.469278  500.465464  340.328969  6.198247  66.712249
1      9123.794425  494.104963  341.748523  7.273029  530.661412

      street      zip      lat      long      unix_time
is_fraud
0      497.153299  496.396598  38.584969 -90.384436  1.358900e+09
1      497.573827  485.901668  38.742813 -90.033730  1.355323e+09

# perform Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(new_df.drop('is_fraud', axis=1), new_df['is_fraud'], test_size=0.3, random_state=41)

# Perform Grid Search Cross-Validation
dt_model = DecisionTreeClassifier()
param_grid = {
    'max_depth': [None, 5, 10, 15],
    'min_samples_split': [5, 10, 15],
    'min_samples_leaf': [2, 4, 8]
}
grid_search = GridSearchCV(estimator=dt_model, param_grid=param_grid, cv=3, scoring='accuracy')
grid_search.fit(X_train, y_train)
best_params = grid_search.best_params_
print(best_params)

{'max_depth': 10, 'min_samples_leaf': 8, 'min_samples_split': 5}

# Create Decision Tree model with the best hyperparameters
best_dt_model = DecisionTreeClassifier(**best_params)
best_dt_model.fit(X_train, y_train)
y_pred = best_dt_model.predict(X_test)

# Evaluate the model
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

Accuracy: 0.961591457113331
Classification Report:
              precision    recall  f1-score   support

     0       0.97       0.95       0.96       2893
     1       0.96       0.97       0.96       2913

 accuracy         0.96         0.96         0.96         5806
 macro avg       0.96         0.96         0.96         5806
 weighted avg    0.96         0.96         0.96         5806


Confusion Matrix:
[[2761  132]
 [  91 2822]]

# Display feature importances
feature_importances = best_dt_model.feature_importances_
print("Feature Importances:")
for feature, importance in zip(X_train.columns, feature_importances):
    print(f"{feature}: {importance:.5f}")

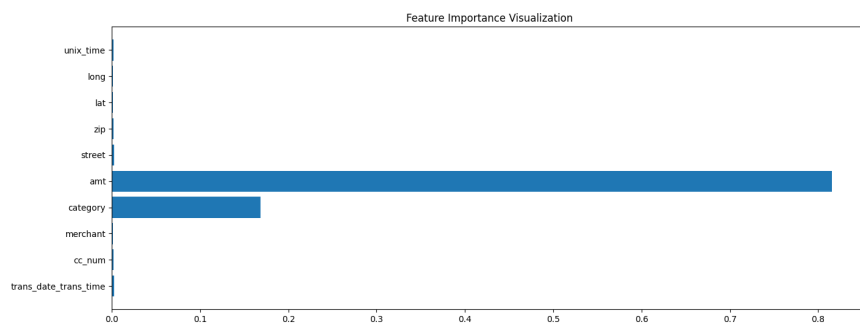
Feature Importances:
trans_date_trans_time: 0.00292
cc_num: 0.00198
merchant: 0.00123
category: 0.16852
amt: 0.81564
street: 0.00274
zip: 0.00216
lat: 0.00137
long: 0.00166
unix_time: 0.00177

#plot correlaion matrix
correlation_matrix=new_df.corr()
plt.figure(figsize=(10,8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title("Correlation Matrix for features")
plt.show()


```




```
#plot feature importances
plt.figure(figsize=(16, 6))
plt.barh(X_train.columns, best_dt_model.feature_importances_)
plt.title('Feature Importance Visualization')
plt.show()
```



```
# Plot the precision-recall curve
precision, recall, _ = precision_recall_curve(y_test, best_dt_model.predict_proba(X_test)[: , 1])
plt.figure(figsize=(8, 6))
plt.plot(recall, precision, color='blue')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.show()
```



```
# Plot the confusion matrix using heatmap
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='g', cbar=False)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
# Visualize the decision tree
plt.figure(figsize=(150, 100))
plot_tree(best_dt_model, filled=True, feature_names=[f"Feature {i}" for i in range(X.shape[1])], class_names=["0", "1"])
plt.show()
```